

▼	xv6	
▼	Chapter 1 Operating system interfaces	9
▼	Introduction	9
•	Shell	10
	The shell is an ordinary program that reads commands from the user and executes them.	
▼	System Calls	9
	When a process needs to invoke a kernel service, it invokes a system call, one of the calls in the operating system’s interface. The system call enters the kernel; the kernel performs the service and returns. Thus a process alternates between executing in user space and kernel space.	
•	kernel space	9
•	user space	9
•	Design Interfaces	9
	The trick in resolving this tension is to design interfaces that rely on a few mechanisms that can be combined to provide much generality.	
▼	1.1 Processes and memory	10
▼	Process	10
	process consists of user-space memory	
•	instructions	10
•	data	10
•	stack	10
•	per-process state private to the kernel	10
•	Switch among processes	10
	Xv6 time-shares processes: it transparently switches the available CPUs among the set of processes waiting to execute.	
▼	System call examples	
▼	fork()	10
	Fork creates a new process	
	child process	
	Fork returns in both the parent and the child	
•	Parent	10
	In the parent, fork returns the child’s PID	

- Child
 - in the child, fork returns zero

10

- wait()
 - Returns the PID of an exited (or killed) child of the current process
 - opies the exit status of the child to the address passed to wait;

11

- exec()
 - The exec system call replaces the calling process’s memory with a new memory image loaded from a file stored in the file system.

12

- Not return to program
 - When exec succeeds, it does not return to the calling program; instead, the instructions loaded from the file start executing at the entry point declared in the ELF header

12

- Shell
 - shell is simple

12

- Why fork and exec are not combined
 - Read command
 - Fork
 - Parent waits
 - Child executes
 - The shell exploits the separation in its implementation of I/O redirection.

- Copy-on-write
 - To avoid the wastefulness of creating a duplicate process and then immediately replacing it (with exec), operating kernels optimize the implementation of fork for this use case by using virtual memory techniques such as copy-on-write (see Section 4.6)

12

- Read command
 - reads a line of input from the user with getcmd

12

- Fork
 - fork, which creates a copy of the shell process

12

- Parent waits
 - The parent calls wait

12

- Child executes
 - while the child runs the command

12

- List of System Call

11

System call	Description
int fork()	Create a process, return child’s PID.
int exit(int status)	Terminate the current process; status reported to wait(). No return.
int wait(int *status)	Wait for a child to exit; exit status in *status; returns child PID.
int kill(int pid)	Terminate process PID. Returns 0, or -1 for error.
int getpid()	Return the current process’s PID.
int sleep(int n)	Pause for n clock ticks.
int exec(char *file, char *argv[])	Load a file and execute it with arguments; only returns if error.
char *sbrk(int n)	Grow process’s memory by n bytes. Returns start of new memory.
int open(char *file, int flags)	Open a file; flags indicate read/write; returns an fd (file descriptor).
int write(int fd, char *buf, int n)	Write n bytes from buf to file descriptor fd; returns n.
int read(int fd, char *buf, int n)	Read n bytes into buf; returns number read; or 0 if end of file.
int close(int fd)	Release open file fd.
int dup(int fd)	Return a new file descriptor referring to the same file as fd.
int pipe(int p[])	Create a pipe, put read/write file descriptors in p[0] and p[1].
int chdir(char *dir)	Change the current directory.
int mkdir(char *dir)	Create a new directory.
int mknod(char *file, int, int)	Create a device file.
int fstat(int fd, struct stat *st)	Place info about an open file into *st.
int stat(char *file, struct stat *st)	Place info about a named file into *st.
int link(char *file1, char *file2)	Create another name (file2) for the file file1.
int unlink(char *file)	Remove a file.

▼ Definition13

A file descriptor is **a small integer** representing **a kernel-managed object** that a process may read from or write to.

By convention, a process reads from file descriptor 0 (standard input), writes output to file descriptor 1 (standard output), and writes error messages to file descriptor 2 (standard error).

• stdin13

file descriptor 0 (standard input)

• stdout13

file descriptor 1 (standard output)

• stderr13

file descriptor 2 (standard error)

▼ How13

obtain a file descriptor

• File, Dir and Device13

opening a file, directory, or device

• Pipe13

creating a pipe

• Duplicate13

duplicating an existing descriptor

▼ read and write13

The read and write system calls read bytes from and write bytes to open files named by file descriptors.

• read13

read(fd, buf, n)

1. reads at most n bytes from the file descriptor fd

2. copies them into buf

• write13

write(fd, buf, n)

1. writes n bytes from buf to the file descriptor fd

2. returns the number of bytes written

• close14

The close system call releases a file descriptor

▼ redirection14

cat < input.txt

- New file descriptor

14

A newly allocated file descriptor is always the lowest numbered unused descriptor of the current process.



```
char *argv[2];

argv[0] = "cat";
argv[1] = 0;
if(fork() == 0) {
    close(0);
    open("input.txt", O_RDONLY);
    exec("cat", argv);
}
```

14

```
case REDIR:
    rcmd = (struct redircmd*)cmd;
    close(rcmd->fd);
    if(open(rcmd->file, rcmd->mode) < 0){
        fprintf(2, "open %s failed\n", rcmd->file);
        exit(1);
    }
    rcmd->fd = open(rcmd->file, rcmd->mode);
    break;
```

- close

14

child closes file descriptor 0

- open

14

0 will be the smallest available file descriptor

- fork and exec are separate calls

14

The shell has a chance to redirect the child's I/O without disturbing the I/O setup of the main shell

▼ hello world into a file

15

- fork

15

```
if(fork() == 0) {
    write(1, "hello ", 6);
    exit(0);
} else {
    wait(0);
    write(1, "world\n", 6);
}
```

▼ dup

15

```
fd = dup(1);
write(1, "hello ", 6);
write(fd, "world\n", 6);
```

dup system call duplicates an existing file descriptor

Both file descriptors share an offset

- Error redirection

| ls existing-file non-existing-file > tmp1 2>&1.

| Both the name

| and the error message

| will show up in the file tmp1

15

▼ 1.3 Pipes

15

- | a small kernel buffer

15

▼ | processes as a pair of file descriptors

15

- | reading

15

- | writing

15

- ```
int p[2];
char *argv[2];

argv[0] = "wc";
argv[1] = 0;

pipe(p);
if(fork() == 0) {
 close(0);
 dup(p[0]);
 close(p[0]);
 close(p[1]);
 exec("/bin/wc", argv);
} else {
 close(p[0]);
 write(p[1], "hello world\n", 12);
 close(p[1]);
}
```

16

- 1.4 File system

17

- 1.5 Real world

19